



Visual Tracking I: Monocular 2D point tracking

Kim Steenstrup Pedersen



Plan for today

- The monocular visual 2D tracking problem
 - Definition of the problem
 - Examples
 - Applications
- The Kanade-Lucas-Tomasi (KLT) feature tracker
 - The basic algorithm
 - Extension: Good features to track



Monocular visual 2D tracking

- Monocular = single camera
- Visual 2D tracking:
 - To follow objects apparent 2 dimensional motion in a video sequence.
 - At every frame, locate objects and associate with objects currently being tracked.
 - One object is “easy” – for multiple objects we need to keep track of the identity of objects.
- Similar to optical flow (dense), but we are interested in motion of regions of the image (sparse).

Example: 2D tracking of object



Results from Hauberg-Lauze-Pedersen, JMIV 2012

Challenges in visual tracking

(Inspired by David J. Fleet)



- Choices to make:
 - What to model / estimate: Shape (2D/3D), appearance, object dynamics.
 - What to measure: Feature points, optical flow, color histograms, edges, etc.
- Some of the main challenges are:
 - Objects with many degrees of freedom, affecting shape, appearance, and motion.
 - Occlusion and large scale changes.
 - Multiple objects and background clutter.



Examples of applications of tracking

- Surveillance (parking lot, airports, shops, etc.)
- Target tracking (military and civil)
- Automated car driver assistance (avoid other cars and pedestrians, stay on the road)
- Device-less human computer interaction (e.g. simple 2D gesture recognition)
- And many more ...



Approaches to 2D visual tracking


- Tracking by visual feature matching:
 - Represent an object by a collection of image features.
 - Tracking is performed by matching features between frames.
- Region / blob tracking:
 - Represent and track the appearance of the interior of the object
- Contour tracking:
 - Represent and track the outline contour of the object

Example: 2D tracking of object (region based)



Results from Hauberg-Lauze-Pedersen, JMIV 2012

Examples of a contour tracker



Isard & Blake., "CONDENSATION" IJCV '98
<http://www.robots.ox.ac.uk/~misard/condensation.html>



Common steps in tracking


- Target representation and localization:
 - Choose an appropriate representation for the target (e.g. position and scale, 2D/3D graphics model, ...)
 - Localization: Estimate the parameters in the target representation.
- Filtering:
 - Noise and numerical instabilities introduce drift in the tracker.
 - This can be handled introducing temporal filtering, such as Bayesian stochastic filtering.
- Data association:
 - How do we make sure that we are tracking the same object in all frames.
 - Especially important for multi-target tracking.

The filtering problem



Results from Hauberg-Lauze-Pedersen, JMIV 2012

Data association is a hard problem: Keeping track of the player IDs



Okuma et al., “Boosted Particle Filter.” ECCV ‘04



Tracking by matching



Kanade-Lucas-Tomasi (KLT) tracker

- Basic Idea:
 - Track features between frames essentially using the Lucas-Kanade optic flow algorithm.
 - Only select good features to track and drop features when they become unstable.
- Video sequence: $I(x, y, t)$
- Local image model: $J(\mathbf{x} + \mathbf{d}) = I(\mathbf{x}) + n(\mathbf{x})$
where $J(\mathbf{x} + \mathbf{d}) = I(x + \xi, y + \eta, t + \tau)$ is the image at $t + \tau$,
 $I(\mathbf{x}) = I(x, y, t)$ is the image at time t ,
the displacement $\mathbf{d} = (\xi, \eta)^T$, and $n(\mathbf{x})$ represent noise.
- We want to find displacement \mathbf{d} that minimize the residue (matching error) on a patch window W :
$$\varepsilon = \int_W (J(\mathbf{x} + \mathbf{d}) - I(\mathbf{x}))^2 w(\mathbf{x}) d\mathbf{x}$$



KLT: Solving for the image displacement

- Linearization of image by truncated Taylor series

$$J(\mathbf{x} + \mathbf{d}) \approx J(\mathbf{x}) + \mathbf{g}^T \mathbf{d}$$

$$\text{where } J(\mathbf{x}) = J(x, y, t + \tau) \text{ , } \mathbf{g} = \left(\frac{\partial J}{\partial x}, \frac{\partial J}{\partial y} \right)^T \text{ , } \mathbf{d} = (\xi, \eta)^T$$

- Plug into residue expression

$$\varepsilon = \int_W \left(J(\mathbf{x}) - I(\mathbf{x}) + \mathbf{g}^T \mathbf{d} \right)^2 w(\mathbf{x}) d\mathbf{x}$$

- A closed form solution exist by differentiation wrt. $\mathbf{d} = (\xi, \eta)^T$

$$\frac{\partial \varepsilon}{\partial \mathbf{d}} = 2 \int_W \left(J(\mathbf{x}) - I(\mathbf{x}) + \mathbf{g}^T \mathbf{d} \right) \mathbf{g} w(\mathbf{x}) d\mathbf{x} = 0$$

Use chain rule
and

$$\frac{d\mathbf{g}^T \mathbf{d}}{d\mathbf{d}} = \mathbf{g}$$

(continues on next slide)



KLT: Solving for the image displacement

$$\frac{d\varepsilon}{d\mathbf{d}} = 2 \int_W \left(J(\mathbf{x}) - I(\mathbf{x}) + \mathbf{g}^T \mathbf{d} \right) \mathbf{g} w(\mathbf{x}) d\mathbf{x} = 0 \Rightarrow$$

• Use that

$$\left(\int_W (\mathbf{g} \mathbf{g}^T) w(\mathbf{x}) d\mathbf{x} \right) \mathbf{d} = - \int_W (J(\mathbf{x}) - I(\mathbf{x})) \mathbf{g} w(\mathbf{x}) d\mathbf{x} \quad (\mathbf{g}^T \mathbf{d}) \mathbf{g} = (\mathbf{g} \mathbf{g}^T) \mathbf{d}$$

- This is just a system of linear equations

$$\mathbf{G} \mathbf{d} = \mathbf{e} \text{ with solution } \mathbf{d} = \mathbf{G}^{-1} \mathbf{e}$$

- 2×2 “Harris” matrix aka structure tensor

$$\mathbf{G} = \int_W (\mathbf{g} \mathbf{g}^T) w(\mathbf{x}) d\mathbf{x} \approx \sum_{\mathbf{x} \in W} \begin{bmatrix} g_x^2(\mathbf{x}) & g_x(\mathbf{x}) g_y(\mathbf{x}) \\ g_x(\mathbf{x}) g_y(\mathbf{x}) & g_y^2(\mathbf{x}) \end{bmatrix} w(\mathbf{x})$$

- Residue projected on gradient

$$\mathbf{e} = \int_W (I(\mathbf{x}) - J(\mathbf{x})) \mathbf{g} w(\mathbf{x}) d\mathbf{x} \approx \sum_{\mathbf{x} \in W} \begin{bmatrix} (I(\mathbf{x}) - J(\mathbf{x})) g_x(\mathbf{x}) w(\mathbf{x}) \\ (I(\mathbf{x}) - J(\mathbf{x})) g_y(\mathbf{x}) w(\mathbf{x}) \end{bmatrix}$$



KLT: Select features for tracking

- Choose a fixed patch size
- We want to only keep patches which provide good displacement estimates – check eigenvalues of **G**:
 - If both eigenvalue are small – nearly flat patch (not good)
 - If one eigenvalue large and one small – edge (not good)
 - Both eigenvalues sufficiently large – corner or texture (good)
- We keep a patch / feature, if

$$\min(\lambda_1, \lambda_2) > \lambda$$


where λ_1, λ_2 are the eigenvalues of **G** and λ a threshold.

- Select non-overlapping patches in image passing the eigenvalue test.



KLT: Example of feature candidates

Patch size = 15 x 15 pixel, eigenvalues threshold = 10



Tomasi & Kanade, Tech. report CMU '91



The KLT algorithm

1. Find new non-overlapping feature patch candidates
 - Check feature stability, keep if $\min(\lambda_1, \lambda_2) > \lambda$
 - Store intensity patch from detection frame and current position of feature patch.
2. For each feature patch
 - (Check feature stability, keep if $\min(\lambda_1, \lambda_2) > \lambda$)
 - Estimate displacement by $\mathbf{d}_{new} = \mathbf{G}^{-1} \mathbf{e}$
 - Update feature position: $\mathbf{x}_{new} = \mathbf{x}_{old} + \mathbf{d}_{new}$
3. Loop back to 1. and process next frame.



The KLT algorithm

1. Find new non-overlapping feature patch candidates
 - Check feature stability, keep if $\min(\lambda_1, \lambda_2) > \lambda$
 - Store intensity patch from detection frame and current position of feature patch.
2. For each feature patch
 - (Check feature stability, keep if $\min(\lambda_1, \lambda_2) > \lambda$)
 - Iterate until convergence in residue $\varepsilon = \int_w (J(\mathbf{x} + \mathbf{d}) - I(\mathbf{x}))^2 w(\mathbf{x}) d\mathbf{x}$
 - Estimate displacement by $\mathbf{d}_{new} = \mathbf{G}^{-1} \mathbf{e}$
 - Resample patch $J(\mathbf{x} + \mathbf{d}_{new})$ with sub-pixel precision by bilinear interpolation.
 - If estimate do not converge – drop feature
 - Update feature position: $\mathbf{x}_{new} = \mathbf{x}_{old} + \mathbf{d}_{new}$
3. Loop back to 1. and process next frame.



Aside: Resampling the patch with sub-pixel precision

- Bilinear interpolation:
See e.g. http://en.wikipedia.org/wiki/Bilinear_interpolation
- Resample patch from J at $J(\mathbf{x} + \mathbf{d}_{new})$ by bilinear interpolation:

Make a drawing on whiteboard

- Use the resampled patch to compute \mathbf{G} , \mathbf{e} and residual ε

Kanade-Lucas-Tomasi (KLT) tracker



GPU_KLT:

A GPU-based Implementation of the
Kanade-Lucas-Tomasi Feature Tracker

Sinha et al. "GPU_KLT: A GPU-based Implementation of the KLT Feature Tracker" 2006
http://cs.unc.edu/~ssinha/Research/GPU_KLT/



KLT: Assumptions

- Linear displacement:
 - Patches must be planar or near planar.
 - High curvature will introduce deformation of the patch violating the linear assumption.
 - As tracking time increases, patches deform in a non-linear fashion or start to straddle occlusion boundaries.
- Patch size:
 - We assume that one patch size fits all – not really correct.
 - Too small patches leads to noisy displacement estimates.
 - Too large patches leads to problems with occlusion boundaries.



Extension: Good features to track

- In order to detect occlusion boundaries, introduce an affine deformation model and a dissimilarity measure.
- Dissimilarity between first patch and current patch:

$$\varepsilon = \int_W \left(J((\mathbf{I} + \mathbf{D})\mathbf{x} + \mathbf{d}) - I(\mathbf{x}) \right)^2 w(\mathbf{x}) d\mathbf{x}$$

- First estimate \mathbf{D} and \mathbf{d} by solving

$$\mathbf{T}\mathbf{z} = \mathbf{a} \Rightarrow \mathbf{z} = \text{pinv}(\mathbf{T})\mathbf{a} = \left(\mathbf{T}^T \mathbf{T} \right)^{-1} \mathbf{T}^T \mathbf{a}$$

where $\mathbf{z} = \left(D_{xx}, D_{xy}, D_{yx}, D_{yy}, d_x, d_y \right)^T$ and see Shi-Tomasi for definition of \mathbf{T} and \mathbf{a} .

- Detect occlusions and instability by thresholding the dissimilarity measure. Remove features that fail this test.

Kanade-Lucas-Tomasi (KLT) tracker



GPU_KLT:


A GPU-based Implementation of the
Kanade-Lucas-Tomasi Feature Tracker

Sinha et al. "GPU_KLT: A GPU-based Implementation of the KLT Feature Tracker" 2006
http://cs.unc.edu/~ssinha/Research/GPU_KLT/



Assignment 4: The last assignment

- Try to implement the KLT tracker.
- Experiment with your tracker on different video sequences.



What's missing?



- I have not covered the **filtering** and **data association** problems.
 - We cover the filtering problem on the *Advanced topics in data modeling* course.
- **Word of advice:** Many tracking algorithms exist, but non of them work in all situation – therefore consider your problem and pick the approach that best meets the challenges in your problem.

Summary



- The monocular visual 2D tracking problem
 - Definition
 - Examples
 - Applications
- The Kanade-Lucas-Tomasi (KLT) feature tracker
 - The basic algorithm
 - Extension from “Good features to track”



Literature

Reading material:

- Tomasi & Kanade: Detection and Tracking of Point Features. CMU Technical report, CMU-CS-91-132, 1991.
- Shi & Tomasi: Good Features to Track. IEEE Proceedings of CVPR' 94, 593 – 600, 1994.

Additional material:

- Isard & Blake: CONDENSATION – Conditional Density Propagation for Visual Tracking. International Journal of Computer Vision, 29(1): 5 – 28, 1998.